# PRAVIN PANDIT

Senior Applied AI Engineer | AI/ML Engineer | GenAI Specialist
pravin.h.pandit@gmail.com | 9657948419
LinkedIn: https://www.linkedin.com/in/pravinpandit | Address: Pune, Maharashtra

## PROFESSIONAL SUMMARY

Senior AI/ML Engineer with 4+ years of experience designing and deploying production-grade AI systems across finance and enterprise domains. Specialized in AI agents, RAG architectures, and end-to-end ML systems, with strong expertise in Python, Machine Learning, and Deep Learning. Proven track record of designing low-latency, scalable AI systems and improving accuracy and performance across enterprise-scale deployments. GenAI certified, with hands-on experience optimizing systems, translating business problems into scalable AI solutions with measurable product impact.

## KEY HIGHLIGHTS

- Delivered scalable AI solutions across finance and enterprise domains at Xoriant and Accenture.
- Certified in Azure AI, Oracle GenAI, Databricks GenAI; finalist at NPCI Bharat AI Quest.
- Architected low - latency AI pipelines, reducing inference time by up to 40% through optimized workflows.
- Built 3+ large-scale GenAI systems including AI Agents, RAG-powered copilots, and Financial Intelligence solutions.
- Improved system accuracy by 40–50% via prompt optimization, tool selection logic, and retrieval tuning.

## TECHNICAL SKILLS

- **Programming & Systems:** Python, Flask
- **Machine Learning:** Supervised & Unsupervised Learning, Feature Engineering, Model Evaluation, Bias–Variance Analysis, PCA, Class Imbalance Handling
- **Deep Learning:** Neural Networks, Transformers (NLP), RNN, LSTM, CNN.
- **Generative AI & LLMs:** AI Agents, RAG Architectures, Prompt Engineering, Tool Calling, Embeddings, LlamaIndex, LangChain, OpenAI & Azure OpenAI, llama.
- **Data & Platforms:** SQL (MySQL), Cosmos DB, Oracle DB, Azure Blob Storage
- **MLOps & Deployment:** MLflow, Airflow, Docker, CI/CD, Cloud Deployment (Azure, AWS)

## WORK EXPERIENCE

### Xoriant, Pune. Jun 2024 – Feb 2026

- Designed a production-grade AI agent framework capable of dynamic tool selection based on user intent and task context, enabling flexible and efficient task handling.
- Architected context-aware tool selection logic, improving system accuracy by 20% and execution efficiency by 30%.
- Built real-time, WebSocket-based inference pipelines enabling low-latency communication between AI agents, workflow engines, and UI layers.
- Designed a dynamic period & scenario extraction system, reducing inference latency by 60% while improving accuracy by 40–50%.
- Led prompt engineering and tool orchestration strategies to ensure consistent, high-quality outputs across diverse enterprise use cases.

### Accenture Solution Pvt Ltd, Pune. Jan 2022 – Jun 2024

- Designed ML pipelines for data ingestion, preprocessing, and model evaluation, improving system performance by 10–20% and used effective data cleaning and data balancing techniques.

- Performed data extraction from Azure Blob Storage, utilizing advanced techniques to refine critical information, achieving a 20% performance enhancement. Managed MySQL and Cosmos database to streamline processing of both offline and online content.
- Leveraged Retrieval-Augmented Generation (RAG) in our system, significantly improving the relevance and accuracy of domain-specific content generation.
- Implemented robust data handling practices to ensure integrity, meet project objectives, and enhance functionality, leading to streamlined processes and improved performance

## PROJECTS

### 1. FNPA Assistant – AI-driven Financial Intelligence Agent June 2024 – Feb 2026

- Developed a production-grade AI agent to dynamically handle company-specific financial queries, improving query resolution efficiency by 20-30%.
- Architected scalable workflows using Python, LlamaIndex, LangChain, and ML libraries, integrating Oracle Database, other technologies, and in-house cloud infrastructure to deliver enterprise-grade performance.
- Engineered dynamic tool selection and advanced prompt optimization strategies, enhancing accuracy by 20% and ensuring consistent output quality.
- Built real-time WebSocket APIs enabling seamless, low-latency communication between the agent and user interface, cutting response times by 30% and driving superior user experience.

### 2. SAP Delivery Copilot – Gen AI System Jan 2022 – May 2024

- Developed a dynamic financial document generator using Generative AI, integrating Azure Functions, Logic Apps, OpenAI LLM, MySQL, Cosmos, Node.js, Python, and Redis.
- Contributed significantly to the implementation of RAG, enhancing content generation for domain-specific documents by 30-40%, significantly boosting system performance.
- Led data retrieved from Azure Blob, applying advanced optimization techniques, achieving a 20% performance gain.
- Established robust data handling procedures to maintain data integrity and align with project objectives, thereby enhancing system functionality.

### 3. Bank Customer Churn Prediction System Jan 2024 – APR 2024

- Built an ML-based classification solution predicting churn with 20–30% performance improvement.
- Applied PCA & ADASYN for enhanced data accuracy and fairness.
- Tuned hyperparameters, improving accuracy by 10–20%.
- Evaluated multiple algorithms, selecting optimal models for high precision, recall, and F1-scores.

## EDUCATIONS

- Sandip Institute of Technology and Research Center, Nashik
  B.E in Information Technology - 2021 8.5 CGPA

## CERTIFICATIONS & AWARDS

- Microsoft Certified: Azure AI Fundamentals
- Databricks - Generative AI Fundamentals
- PCAP™ - Certified Associate Python Programmer
- Accenture Tech leap – AI ML Foundation